




Yuxuan Jiang

✉ jyuxuan@umich.edu  essoz.github.io  essoz  Google Scholar

Education

University of Michigan

Sep 2023 to present

Ph.D. in Computer Science and Engineering

- **Research Areas:** Reliable Machine Learning (enabled) Systems, AIOps, Cloud Reliability
- **Advisor:** Prof. Ryan Huang

University of Illinois Urbana-Champaign

Sep 2019 to May 2023

B.E. in Electrical and Computer Engineering

Zhejiang University

Sep 2019 to May 2023

B.E. in Electrical and Computer Engineering

Research Interest

Computer Systems: Enhancing systems reliability and optimizing performance and flexibility through innovative abstraction-level strategies.

Research Experience

Research Assistant

Ann Arbor, MI, USA

University of Michigan – OrderLab, Advised by Prof. Ryan Huang

Aug 2023 to present

- Keywords: **ML-enabled Systems Reliability, Distributed Systems Fail-Slowness Tolerance**
- Leading a project on **detecting silent errors in ML training workflows**—this work addresses a critical gap in the reliability of ML systems by automating checks to prevent errors that degrade model quality. Currently in submission to OSDI'25.
- Conducted a comprehensive analysis of over **50** silent issues across ML pipelines, revealing insights into issues that often go unnoticed and impact model outcomes in real-world settings.
- Developed a monkey-patching-based Python-level instrumentor with **mild overhead (<0.5x) under full instrumentation**. It captures both API invocations and variable updates to provide a **high-fidelity trace** for ML invariant inference.
- Designed a robust analysis pipeline that extracts and refines invariants from program traces. **Precondition inference as its core ensures invariants are accurate, transferable, and diagnostic in nature.**
- Validated the tool's effectiveness in real-world scenarios by successfully detecting over **15 diverse training issues** using invariants generated from **5 small-scale PyTorch official tutorial pipelines**.

Cloud Reliability Research Intern

Beijing, China

Microsoft Research Asia – DKI Group, advised by Dr. Shilin He

Oct 2022 to Jul 2023

- Keywords: **AIOps, ML4Sys, Failure Diagnosis**
- Leveraging LLMs and finetuned models to **automate cloud incident diagnosis** on Azure. This research, recognized at **ICSE'24**, tackled key pain points in cloud operations by generating precise, actionable diagnosis queries from incident tickets.
- Innovated with **curriculum learning and Retrieval-Augmented Generation (RAG)** to enable accurate generation of diagnosis queries derived from incident tickets.
- Created a novel **code-semantics-based** metric for evaluating query generation systems. By statically canonicalizing queries and using **symbolic and dataflow analysis**, this metric allows for accurate comparison **without actual execution**—outperforming traditional NLP metrics.
- Collaborated with product teams to push for real-world deployment of the query generation tool.
- Collaborated with the systems team on data-driven fault detection for large GPU clusters, a project later accepted at ATC'24, advancing reliability for large-scale hardware deployments.

Research Assistant

Urbana, IL

- Keywords: **Controller Reliability, Distributed Systems Fail-Slowness Tolerance**
- Contributed to an **automated testing tool for Kubernetes controllers correctness**, a project later accepted to **SOSP'23**.
- Designed an advanced input exploration strategy to **rigorously test the error-handling pathways of operators**, creating oracles that identify unhealthy system states. This enriched alarm descriptions and reduced false negatives, greatly improving system reliability.
- Interfaced with Kubernetes' ApiMachinery to capture system snapshots, streamlining debugging by allowing for rich state-based diagnostics.
- Re-architected the testing framework (Acto) to enable **parallel testing** and support **multiple runtime environments** like Kind, K3D, and MiniKube.
- Validated Acto's effectiveness on real-world Kubernetes controllers, discovering **18** previously unknown bugs (32% of all identified by Acto) across three popular operators. To date, 14 of these bugs have been patched.
- Worked on a fault injection tool to test distributed systems' tolerance of hardware slowness faults.

Publications

One-Size-Fits-None: Understanding and Enhancing Slow Fault Tolerance in Modern Distributed Systems	Apr 2025
Ruiming Lu, Yunchi Lu, Yuxuan Jiang , Guangtao Xue, Peng Huang	
<i>NSDI'25: Proceedings of the 22nd USENIX Symposium on Networked Systems Design and Implementation</i>	
Xpert: Empowering Incident Management with Query Recommendations via Large Language Models	Apr 2024
Yuxuan Jiang , Chaoyun Zhang, Shilin He, Zhihao Yang, Minghua Ma, Si Qin, Yu Kang, Yingnong Dang, Saravan Rajmohan, Qingwei Lin, Dongmei Zhang	
<i>ICSE'24: Proceedings of the IEEE/ACM 46th International Conference on Software Engineering</i>	
Acto: Automatic End-to-End Testing for Operation Correctness of Cloud System Management	Oct 2023
Jiawei Tyler Gu, Xudong Sun, Wentao Zhang, Yuxuan Jiang , Chen Wang, Mandana Vaziri, Owolabi Legunsen, Tianyin Xu	
<i>SOSP'23: Proceedings of the 29th Symposium on Operating Systems Principles</i>	

Awards

Award of Excellence: Top 5% interns. Microsoft Research Asia (MSRA) 2023
Scholarship for Outstanding Students: Zhejiang University, 2020, 2021, 2022

Leadership & Teaching Experience

Microsoft Student Learn Ambassador Zhejiang University	Nov 2021 to May 2023
Teaching Assistant for CS428: Distributed Systems Zhejiang University	Jan 2023 to May 2023
Academic and Life Tutor for International Students in China Zhejiang University	Oct 2021 to Jan 2022